Gate-initialisation study · GLA, 5 layers, 2.67M params · zero gate weight unless legacy · same data, schedule and lr in every run

A · γ₀ = 0.5 learns recall far better than 0.88 or 0.9975

pass 1 · lines = mean over seeds 0–2, bands = min–max

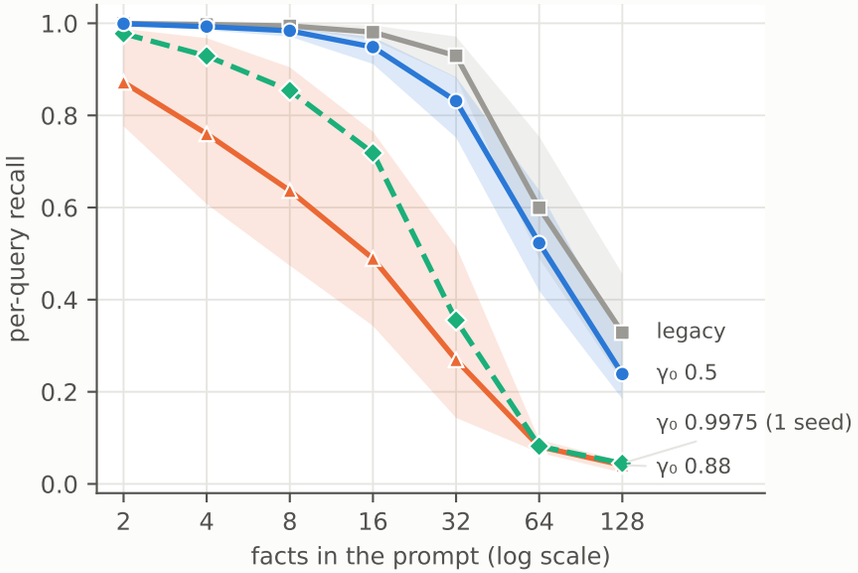

B · γ₀ = 0.5 learns faster

full training record per run · dot = evaluated checkpoint

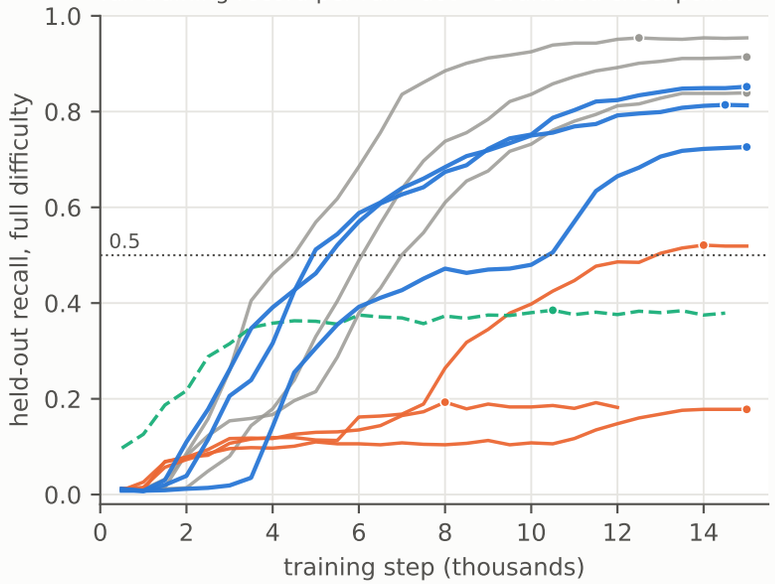

C · Working runs forget fast in layer 0

one point per run · passes 8 and 1

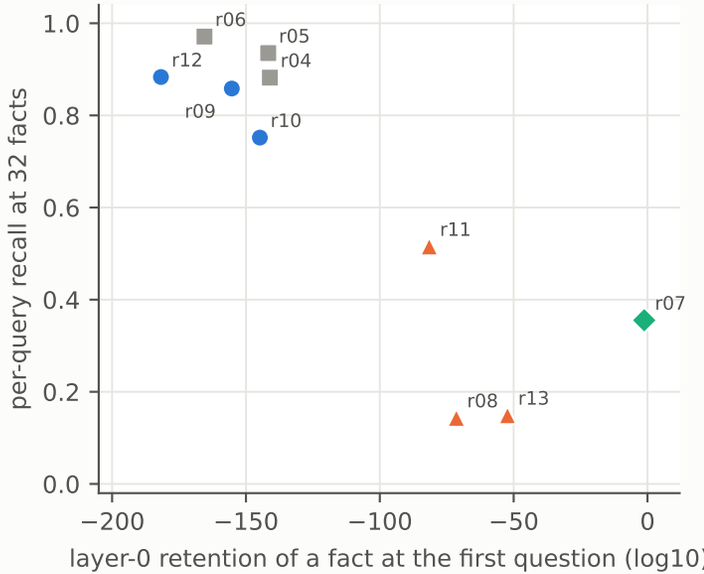

γ₀ = 0.5 (bias 0) · γ₀ = 0.88 (bias 2) · γ₀ = 0.9975 (bias 6) · legacy (bias 0, random W)